



SIMATS
ENGINEERING



SIMATS
Saveetha Institute of Medical And Technical Sciences
(Declared as Deemed to be University under Section 3 of UGC Act 1956)

Knowledge-Based Navigation of an Autonomous Mobile Robot (AMR) PROJECT REPORT

A Project Report submitted in partial fulfilment of the requirements for the Course of

CSA1712 – ARTIFICIAL INTELLIGENCE

to the award of the degree of

BACHELOR OF ENGINEERING

IN

COMPUTER SCIENCE AND ENGINEERING

Submitted by

YUVASRI.S

Under the Supervision of

Dr. R. Veeramani

**Saveetha Institute of Medical and Technical Sciences
Chennai-602105**

SEPTEMBER 2026

Intelligent Knowledge-Based Navigation of an Autonomous Mobile Robot in a Manufacturing Environment

1. Observation and Knowledge Representation

An Intelligent Autonomous Mobile Robot (AMR) operating in a manufacturing plant continuously receives data from cameras, LiDAR, ultrasonic sensors, proximity sensors, wheel encoders, and other sensing devices. However, raw sensor observations alone are not sufficient for intelligent behaviour. The robot must interpret these observations, convert them into structured knowledge, and use that knowledge to make safe and efficient decisions.

Difference Between Observation, Information, and Knowledge

Raw observation is the direct data obtained from sensors. For example, a LiDAR sensor may detect an object 1.5 metres in front of the robot, or a camera may detect a human in the robot's path.

Interpreted information is obtained when the raw observation is processed and given meaning. For example, the robot may interpret the LiDAR reading as "an obstacle is present in the current navigation path."

Useful knowledge represents relationships, rules, and previously learned information that can be used for reasoning and decision-making. For example:

- IF an obstacle blocks the current route, THEN stop or select an alternative route.
- IF a path width is below the safe threshold, THEN reduce speed.
- IF a human is detected within the safety zone, THEN stop or maintain a very low speed.
- IF the route is clear and sufficiently wide, THEN the robot may increase its speed within the permitted limit.

Thus, the transformation can be represented as:

Sensor Data → Interpretation → Knowledge → Reasoning → Decision → Action

Representation of Environmental Observations

The robot can represent environmental knowledge using facts, rules, maps, semantic labels, probability values, and learned models.

For example:

Observation	Interpretation	Knowledge	Possible Action
Object detected ahead	Route is blocked	Current path is unsafe	Stop/avoid obstacle
Small distance between machines	Path is narrow	Reduced clearance exists	Reduce speed
Wide clear corridor	Path is open	Route is suitable	Increase speed safely
Human detected	Human is nearby	High safety risk	Slow down/stop
Forklift approaching	Moving obstacle	Collision risk	Stop/wait/re-route
Restricted zone detected	Area is unavailable	Route cannot be used	Select another route

Industrial Situations

Pallet blocking the normal route:

The robot identifies the pallet using LiDAR or cameras and determines that the normal route is blocked. It can then stop, calculate an alternative path, and continue only after confirming that the alternative route is safe.

Narrow passage between two machines:

The robot estimates the available clearance using sensor measurements. If the clearance is close to the minimum safe width, it reduces its speed and increases its caution while passing through the area.

Temporarily restricted production area:

If the robot receives information that a production area is temporarily restricted, this information can be represented as a temporary "no-entry" zone in its map. The navigation system then avoids that region.

Clear and wide transportation corridor:

If multiple sensors confirm that the corridor is clear and sufficiently wide, the robot can classify it as a preferred route and travel at an appropriate higher speed while remaining within safety limits.

Moving workers or forklifts:

A moving object should not be treated like a stationary obstacle. The robot should estimate its direction and speed and predict whether its future path will intersect with the robot's route. If the collision risk is high, the robot should slow down, stop, or change direction.

2. Knowledge-Based Decision Making

Knowledge-based decision-making allows the AMR to convert environmental conditions into appropriate navigation actions. Instead of responding to individual sensor readings independently, the robot combines observations with stored rules, maps, safety constraints, and previous experience.

Examples of decision rules include:

- **IF** an obstacle blocks the current route, **THEN** stop and search for an alternative route.
- **IF** the path is narrow, **THEN** reduce speed.
- **IF** the path is open and safe, **THEN** increase speed within the permitted limit.
- **IF** a human is detected nearby, **THEN** reduce speed or stop.
- **IF** multiple obstacles block the route, **THEN** calculate another route.
- **IF** sensor information is uncertain, **THEN** choose the safer action or stop until the environment becomes clear.

Combining Multiple Observations

Individual observations may not provide sufficient information. Combining multiple observations creates more reliable knowledge.

For example:

- LiDAR detects an object.
- Camera identifies the object as a worker.
- Proximity sensor confirms that the worker is close.
- Tracking information shows that the worker is moving toward the robot's path.

Together, these observations produce the knowledge:

"A moving human is approaching the robot's operating area, creating a high collision risk."

The appropriate decision would therefore be to reduce speed or stop.

Similarly, if LiDAR shows an apparently narrow path but the camera and additional sensors indicate that part of the apparent obstruction is a harmless reflection or temporary sensor error, the robot can reassess the situation before making a navigation decision.

This demonstrates that **sensor fusion** provides more reliable knowledge than depending on a single sensor.

3. Transformation of Observations into Knowledge

Knowledge formation is not a one-time process. The robot continuously observes its environment and updates its internal representation.

The general process is:

Observation → Interpretation → Knowledge → Decision

For example:

1. **Observation:** An object is detected in front of the robot.
2. **Interpretation:** The object appears to be a pallet.
3. **Knowledge:** The normal route is temporarily blocked.
4. **Decision:** Stop and select an alternative route.
5. **Action:** Navigate around the pallet.

Repeated observations allow the robot to develop more general knowledge and improve future decisions.

Temporary vs. Permanent Obstacles

The robot can compare the current observation with its stored map and previous observations. If an obstacle appears suddenly and disappears later, it can be classified as temporary. If the object remains at the same location for a long period and matches the plant map, it may be classified as permanent.

Genuine Narrow Path vs. False Detection

A narrow-path classification should ideally be confirmed using multiple sensors and repeated measurements. If only one sensor reports a narrow path while other sensors show sufficient clearance, the robot should assign lower confidence to the observation and investigate further.

Completely Open vs. Partially Obstructed Path

A path should not be classified as completely open merely because its centre is clear. The robot should consider the entire required safety corridor around its body. Partial obstructions near the edges can still create collision risks, particularly when carrying wide loads.

Stationary vs. Moving Object

The robot can compare successive observations to determine whether an object changes position. Object tracking can provide information about movement direction and speed. This helps distinguish a stationary pallet from a moving forklift or worker.

Safe vs. Uncertain Environment

A safe environment is one where the robot has sufficient reliable information about the route, clearance, obstacles, and surrounding movement. If important information is missing or sensors disagree significantly, the environment should be classified as uncertain. The robot should then adopt a conservative action.

4. Handling Noisy Observations

Industrial environments are challenging because sensor readings can be affected by dust, poor lighting, reflections, vibrations, machinery movement, electromagnetic interference, sensor limitations, and communication failures.

Effects of Noisy Observations

If an obstacle is detected when no obstacle exists, the robot may unnecessarily stop or change its route. This can reduce productivity and cause delays.

If an actual obstacle is temporarily missed, the consequences can be much more serious because the robot may continue moving toward a collision.

If a narrow path is incorrectly classified as open, the robot may travel too quickly through an unsafe area.

If different sensors provide contradictory readings, the robot may be unable to determine the correct environmental state.

For example:

- LiDAR: obstacle detected.
- Camera: path appears clear.
- Ultrasonic sensor: object detected at close range.

Instead of immediately trusting one sensor, the robot should perform **sensor fusion**, assess sensor reliability, and consider the uncertainty associated with each observation.

Methods for Handling Noise

The robot can use:

- Sensor filtering.
- Multiple sensor measurements.
- Sensor fusion.
- Confidence scores.
- Repeated observations.
- Object tracking.
- Probabilistic reasoning.

- Outlier detection.
- Safety margins.

For example, if one sensor briefly detects an obstacle but subsequent measurements from several sensors show no obstacle, the robot can classify the original reading as potentially noisy. However, safety-critical observations should still be treated conservatively until sufficient evidence is available.

Therefore, relying on a single observation can result in inappropriate or unsafe decisions.

Reliable navigation requires multiple observations and confidence-based reasoning.

5. Handling Incomplete Knowledge

The robot may not always have complete information about its environment. Sensors have limited range and may be blocked by machines, pallets, walls, or other objects.

Examples of Incomplete Information

Path outside sensor range:

The robot may know that the current section of the corridor is clear but may have no information about what is beyond the sensing range. It should therefore avoid assuming that the entire route is safe.

Forklift suddenly entering the route:

A forklift may enter the robot's path from a side corridor. The robot must continuously update its environment model and react to newly detected information.

Worker hidden behind a machine:

The robot may not detect a worker until the worker becomes visible. Therefore, areas with limited visibility should be treated with additional caution.

Temporary sensor failure:

If a camera or LiDAR unit becomes unavailable, the robot should determine whether the remaining sensors provide sufficient information for safe operation. If not, it should reduce speed or stop.

Insufficient information about route safety:

When the robot cannot confidently determine whether a route is safe, it should not assume

that the route is safe. Instead, it should choose a safer route, slow down, stop, or gather additional information.

Importance of Uncertainty

Incomplete information introduces uncertainty into decision-making. An intelligent robot should therefore distinguish between:

- **Known safe**
- **Known unsafe**
- **Unknown/uncertain**

This is important because treating an unknown condition as safe can lead to collisions. A safety-oriented robot should generally adopt a conservative action when uncertainty is high.

6. Conflicting and Ambiguous Observations

Different sensors may provide different interpretations of the same environment. This creates conflicting or ambiguous observations.

For example:

- LiDAR reports an obstacle while the camera reports an open path.
- The camera identifies a wide corridor while LiDAR detects a partial obstruction.
- A sensor detects an object but cannot determine whether it is a stationary machine component or a moving worker.

Causes of Conflicting Observations

Conflicts may occur because:

- Sensors have different operating principles.
- Sensors have different fields of view.
- Lighting affects camera performance.

- Reflective surfaces affect LiDAR.
- Ultrasonic sensors may have limited resolution.
- Objects may be partially hidden.
- Sensors may be temporarily malfunctioning.
- Measurements may be taken at slightly different times.

Resolving Conflicts

The robot can use sensor fusion and confidence values to combine information.

For example:

Camera: 70% confidence that the path is clear.

LiDAR: 95% confidence that an object is present.

Ultrasonic: 90% confidence that an object is nearby.

The robot should give greater importance to the stronger evidence, especially when safety is involved. It should also consider the type of sensor and its suitability for detecting the particular object.

If the robot cannot resolve the conflict with sufficient confidence, it should classify the situation as uncertain and adopt a safe action such as slowing down or stopping.

Ambiguous objects should also be treated conservatively. If the robot cannot determine whether an object is a harmless stationary component or a moving worker, it should not assume that it is harmless.

7. Safety and Industrial Constraints

Safety must have a higher priority than speed and productivity in an autonomous industrial robot. A small reduction in productivity is preferable to a collision involving workers, equipment, materials, or the robot itself.

The robot's knowledge base should therefore include safety constraints such as:

- Minimum obstacle clearance.

- Maximum permitted speed.
- Human safety zones.
- Emergency stopping conditions.
- Minimum turning radius.
- Load limitations.
- Restricted areas.
- Minimum visibility requirements.
- Safe stopping distance.

Appropriate Responses

Sudden obstacle:

The robot should immediately reduce speed or stop and then determine whether an alternative route is available.

Human enters operating area:

The robot should reduce speed or stop depending on the distance, predicted movement, and safety rules.

Uncertain path:

The robot should select the safer action rather than assuming the path is clear.

Multiple routes:

The robot can compare routes based on safety, distance, congestion, travel time, and energy consumption. Safety constraints should be applied first.

Higher speed:

Although increased speed may improve productivity, it also reduces available reaction time and increases stopping distance. Therefore, speed should only be increased when environmental conditions and safety margins permit it.

The robot's objective can therefore be considered as:

Safety → Feasibility → Reliability → Efficiency

rather than simply attempting to minimize travel time.

8. Knowledge Updating

Manufacturing environments are dynamic. Machines may be moved, production areas may change, workers may enter or leave, and temporary obstacles may appear. Consequently, the robot's knowledge must be continuously updated.

New observations can perform several functions:

Confirm Existing Knowledge

Repeatedly observing a clear corridor can increase confidence that the corridor is normally available.

Correct Existing Knowledge

If the robot's map indicates that a corridor is open but repeated observations show that a new machine has been installed, the robot should update its map.

Add New Knowledge

If a new restricted production zone is identified, the robot can add the restricted area to its knowledge base.

Remove Outdated Knowledge

If a temporary pallet blocking a route is removed, the robot should remove the corresponding temporary obstacle information.

Change Current Decision

If the robot is travelling toward a destination and a forklift suddenly blocks the route, the newly acquired information should override the previous navigation decision.

This continuous updating allows the robot to remain responsive to changes rather than relying on outdated information.

9. Overall Analysis

An intelligent industrial AMR can transform simple sensor observations into intelligent autonomous behaviour through a continuous knowledge-based process:

Sensor Observation → Interpretation → Knowledge Formation → Reasoning → Decision → Action

For example:

Sensor Observation:

LiDAR detects an object 2 metres ahead.

↓

Interpretation:

The object is identified as a pallet occupying the robot's path.

↓

Knowledge Formation:

The current route is temporarily blocked.

↓

Reasoning:

The robot checks alternative routes and safety constraints.

↓

Decision:

An alternative route is selected.

↓

Action:

The robot slows down, turns safely, and follows the alternative route.

If a worker subsequently enters the alternative route, the robot updates its knowledge and changes its decision again.

Benefits of the Knowledge-Based Approach

1. Improved Navigation Safety:

The robot can reason about obstacles, humans, restricted zones, and uncertain conditions instead of reacting blindly to sensor readings.

2. Improved Operational Efficiency:

When the environment is confirmed to be safe, the robot can select efficient routes and increase speed within safety limits.

3. Better Collision Avoidance:

Combining sensor data and object tracking allows the robot to detect stationary and moving obstacles and predict potential collisions.

4. Adaptability:

Continuous knowledge updating allows the robot to respond to temporary obstacles, changing production layouts, restricted areas, and moving objects.

5. Decision-Making Under Uncertainty:

Confidence values, sensor fusion, repeated observations, and conservative safety rules allow the robot to make safer decisions when information is incomplete or conflicting.

6. Autonomous Operation:

The robot can independently observe its environment, form knowledge, reason about possible actions, and execute appropriate navigation decisions without requiring continuous human intervention.

Conclusion

An autonomous mobile robot in a manufacturing environment must do more than collect sensor readings. It must transform raw observations into interpreted information and structured knowledge that can support reasoning and decision-making. Observations such as obstacles, narrow paths, open corridors, workers, forklifts, and restricted areas become useful when they are combined with rules, maps, previous observations, safety constraints, and confidence information.

Because industrial environments contain noise, uncertainty, incomplete information, conflicting sensor readings, and continuously changing conditions, the robot should not rely on isolated observations. Sensor fusion, repeated measurements, probabilistic reasoning,

object tracking, safety margins, and continuous knowledge updating are essential for reliable operation.

The complete intelligent navigation process can therefore be represented as:

Observation → Interpretation → Knowledge → Reasoning → Decision → Action → Knowledge Update

This knowledge-based approach enables the AMR to achieve safer navigation, better collision avoidance, improved operational efficiency, adaptability to changing environments, and reliable autonomous operation. Most importantly, when safety and productivity conflict, the robot's knowledge and reasoning system should always prioritize **human and environmental safety over speed or efficiency**.